

# call/cc in Ten Minutes

Spencer Tipping

February 24, 2012

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>return</b>	<b>2</b>

## 1 Introduction

I'm going to go off on an inexplicable tangent for a minute here, but I promise it has something to do with continuations.

I was writing an ext4 filesystem driver recently and came across something I didn't expect. You know how you can say `cd ..` to go up one directory? I always assumed this worked by removing a directory from your current location, sort of like returning from a function. But it doesn't work this way at all. Instead, every directory in ext4 actually has a *subdirectory* called `..` that you can *cd into*. This special subdirectory points back to the parent of whichever directory you're in.<sup>1</sup>

At first this made no sense to me because it meant that your path would just keep growing forever. But then I realized that programs don't store their path, they store only their current directory (by its inode, or disk-level data structure that represents it). This meant that it didn't matter how you got to a directory, it only mattered where you were and where you could go from there. (The path can be constructed after the fact by continuously *cding into ..* until the root is reached. Wild, huh?)

Now imagine that your path is like a function call stack. You *cd* into a subdirectory, you do stuff, and you *cd* back out of it. Logically this resembles a single function call, and if `cd ..` were implemented as an "undo" of the last *cd* it would indeed be a real stack. However, each *cd* is doing the same thing: It's calling into a subdirectory of wherever you are. And this is exactly how continuations work.

---

<sup>1</sup>Each directory also has a subdirectory called `.`, which points back to the current directory.

## 2 return

If you've written code in any normal programming language like C, Java, Javascript,<sup>2</sup> etc, you've probably written the word `return`. And further, you've probably used it to escape from a function earlier than you normally would. For instance:

```
var factorial = function (n) {  
  if (n === 0) return 1;  
  return n * factorial(n - 1);  
};
```

If the first `return` is hit, the second one will never happen. So `return` is really doing two things; it's setting up the value that the function will have, and it's jumping back to the parent function. The jump works because when we called into `factorial`, we left the return address on the stack. `return` doesn't really know where to go, it just goes to whichever address the caller specifies.

Let's step back for a moment and think about a different interpretation of the word `return`. What if it were a function? The call into `factorial` is a jump, and the return from `factorial` is also a jump (just like `cd directory` and `cd ..` are both jumps). Are they really the same thing?

As it turns out, they are indeed.<sup>3</sup> Here's the normal way you'd call `factorial`:

```
alert('About to call factorial');  
var x = factorial(5);  
alert('The factorial of 5 is ' + x);  
exit();    // Exits immediately and never returns.
```

And here's what it looks like when we model `return` as a function:

```
var factorial = function (n, ret) {  
  if (n === 0) ret(1);  
  ret(n * factorial(n - 1));  
};  
  
alert('About to call factorial');  
factorial(5, function (value) {  
  var x = value;  
  alert('The factorial of 5 is ' + x);  
  exit();    // Exits immediately and never returns.  
});
```

*Todo:* explain the `exit()` function.

---

<sup>2</sup>Ok, Javascript isn't normal, but you get the idea.

<sup>3</sup>With the exception of how state is saved, which I'll get to in a bit.